

# Truncated SVD- Amazon fine food reviews

This Kernel explores the TruncatedSVD and how it can be used on Amazon Food reviews to gather meaningful Clusters

```
[49]: import os
      print(os.listdir("../input"))
```

```
['database.sqlite', 'hashes.txt', 'Reviews.csv']
```

```
[ ]:
```

```
[ ]: # Import all the required libraries

%matplotlib inline
import warnings
warnings.filterwarnings("ignore")

import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn import model_selection
from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_val_score
```

```

12/13 from sklearn.metrics import accuracy_score
      from sklearn.model_selection import RandomizedSearchCV
      from sklearn.model_selection import GridSearchCV

      from collections import Counter

      from sklearn import metrics
      from sklearn.metrics import accuracy_score
      from sklearn.metrics import confusion_matrix
      from sklearn.metrics import roc_curve, auc
      from sklearn.metrics import classification_report
      from sklearn.metrics import average_precision_score

      # Tutorial about Python regular expressions: https://pymotw.com/2/re/
      import re
      import string
      import nltk.corpus
      from nltk.corpus import stopwords
      from nltk.stem import PorterStemmer
      from nltk.stem.porter import PorterStemmer
      from nltk.stem.wordnet import WordNetLemmatizer

      from gensim.models import Word2Vec
      from gensim.models import KeyedVectors
      import pickle

```

```

[ ]: # using the SQLite Table to read data.
      import sqlite3
      show_tables = "select tbl_name from sqlite_master where type = 'table'"
      conn = sqlite3.connect('../input/database.sqlite')
      pd.read_sql(show_tables,conn)

```

```

[ ]: # Data cleaning steps
      #filtering only positive and negative reviews i.e. SKIPPING reviews with Score=3

      filtered_data = pd.read_sql_query("""SELECT * FROM Reviews WHERE Score != 3""")

      # Give reviews with Score>3 a positive rating, and reviews with a score<3 a neg

```

```

12/13 def partition(x):
        if x < 3:
            return '0'
        return '1'

#changing reviews with score less than 3 to be positive and vice-versa
actualScore = filtered_data['Score']
positiveNegative = actualScore.map(partition)
filtered_data['Score'] = positiveNegative

#Display attributes
filtered_data.head()

```

```

[ ]: #Data Cleaning: Deduplication:
display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND UserId="AR5J8UI46CURR"
ORDER BY ProductID
""", conn)
display.head()

#Sorting data according to ProductId in ascending order
sorted_data=filtered_data.sort_values('ProductId', axis=0, ascending=True, inp

#Deduplication of entries
final=sorted_data.drop_duplicates(subset={"UserId","ProfileName","Time","Text"}
final.shape

#Checking to see how much % of data still remains
(final['Id'].size*1.0)/(filtered_data['Id'].size*1.0)*100

# Cleadning the data for Helpfulness duplication

display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND Id=44737 OR Id=64422
ORDER BY ProductID
""", conn)

```

```

12/13 final=final[final.HelpfulnessNumerator<=final.HelpfulnessDenominator]

#Before starting the next phase of preprocessing Lets see the number of entries
print(final.shape)

#How many positive and negative reviews are present in our dataset?

```

```

[ ]: # Get 10k Pts for this analysis

from random import sample
final_dataset = final.ix[np.random.choice(final.index, 10000)]

```

```

[ ]: #Sorting data according to Time in ascending order
KMEANS_DATASET=final_dataset.sort_values('Time', axis=0, ascending=True, inplace=True)

```

```

[ ]: # find sentences containing HTML tags (DATASET FOR BRUTEFORCE)
import re
i=0;
for sent in KMEANS_DATASET['Text'].values:
    if (len(re.findall('<.*?>', sent))):
        print(i)
        print(sent)
        break;
    i += 1;

# find sentences containing HTML tags (DATASET FOR KD_TREE)
import re
i=0;
for sent in KMEANS_DATASET['Text'].values:
    if (len(re.findall('<.*?>', sent))):
        print(i)
        print(sent)
        break;
    i += 1;

# STEMMING
import nltk.corpus
from nltk.corpus import stopwords

```

```

12/13 stop = set(stopwords.words('english')) #set of stopwords
sno = nltk.stem.SnowballStemmer('english') #initialising the snowball stemmer

def cleanhtml(sentence): #function to clean the word of any html-tags
    cleanr = re.compile('<.*?>')
    cleantext = re.sub(cleanr, ' ', sentence)
    return cleantext

def cleanpunc(sentence): #function to clean the word of any punctuation or special characters
    cleaned = re.sub(r'[?|!|\'|"|#]', r'', sentence)
    cleaned = re.sub(r'[.,|)|(|\\|/]', r' ', cleaned)
    return cleaned

#print(stop)
#print('*****')
#print(sno.stem('tasty'))

```

```

[ ]: # LEMMATIZATION
i=0
str1=' '
final_string=[]
all_positive_words=[] # store words from +ve reviews here
all_negative_words=[] # store words from -ve reviews here.
s=''

for sent in KMEANS_DATASET['Text'].values:
    filtered_sentence=[]
    #print(sent);
    sent=cleanhtml(sent) # remove HTML tags
    for w in sent.split():
        for cleaned_words in cleanpunc(w).split():
            if((cleaned_words.isalpha()) & (len(cleaned_words)>2)):
                if(cleaned_words.lower() not in stop):
                    s=(sno.stem(cleaned_words.lower()).encode('utf8'))
                    filtered_sentence.append(s)
                    if (KMEANS_DATASET['Score'].values)[i] == '1':
                        all_positive_words.append(s) #List of all words used to
                    if(KMEANS_DATASET['Score'].values)[i] == '0':
                        all_negative_words.append(s) #List of all words used to
            else:
                continue
        else:
            continue
    else:
        continue

```

```

12/13/      #print(filtered_sentence)
          str1 = b" ".join(filtered_sentence) #final string of cleaned words
          #print("*****")

          final_string.append(str1)
          i = i + 1

```

```

[ ]:      KMEANS_DATASET['CleanedText']=final_string #adding a column of CleanedText which is the cleaned words
          KMEANS_DATASET['CleanedText']=KMEANS_DATASET['CleanedText'].str.decode("utf-8")

```

```

[ ]:      # define column names
          names = ['CleanedText']

          # create design matrix X and target vector y
          X = KMEANS_DATASET[names]
          Y = KMEANS_DATASET['Score']

          X_train, X_test = model_selection.train_test_split(X, test_size=0.1, random_state=42)

```

## Getting the Top 2000 TFIDF features

```

[ ]:      tf_idf_vect = TfidfVectorizer()

          tfidf = tf_idf_vect.fit(X['CleanedText'].values)

          tfidf_train = tfidf.transform(X['CleanedText'].values)

          #tfidf_test = tfidf.transform(X_test['CleanedText'].values)

          features = tf_idf_vect.get_feature_names()

```

```

[ ]:      features_set= (np.asarray(features)[0:5,])
          print(features_set)

```

# MODULE 1: ANALYSIS OF TOP 2000 TFIDF FEATURES

## Part 1: get the Top 2000 TFIDF Features

```
[ ]: # source: https://buhrmann.github.io/tfidf-analysis.html
def top_tfidf_feats(row, features, top_n=10000):
    ''' Get top n tfidf values in row and return them with their corresponding
    topn_ids = np.argsort(row)[:top_n]
    #top_feats = [(features[i], row[i]) for i in topn_ids]
    top_feats = [(features[i]) for i in topn_ids]
    df = pd.DataFrame(top_feats)
    #df.columns = ['feature', 'tfidf']
    df.columns = ['feature']
    return top_feats

top_tfidf = top_tfidf_feats(tfidf_train[1,:].toarray()[0], features, 2000)
```

## Part 2 : Created Word Co-Occurrence Matrix with neighbourhood = 5

```
[ ]: import numpy as np
import pandas as pd

ctxs = list(top_tfidf)

l_unique = list(set((' '.join(ctxs)).split(' ')))
mat = np.zeros((len(l_unique), len(l_unique)))

nei = []
nei_size = 5

for ctx in ctxs:
    words = ctx.split(' ')

    for i, _ in enumerate(words):
```

https://

12/13/

```

nei.append(words[i])

if len(nei) > (nei_size * 2) + 1:
    nei.pop(0)

pos = int(len(nei) / 2)
for j, _ in enumerate(nei):
    if nei[j] in l_unique and words[i] in l_unique:
        mat[l_unique.index(nei[j]), l_unique.index(words[i])] += 1

mat = pd.DataFrame(mat)
mat.index = l_unique
mat.columns = l_unique
print(mat)

```

## Part 3: Word Co-Occurrence Matrix Decomposition using SVD

```

[ ]: import numpy as np

la= np.linalg

#u, s, vh = la.svd(mat, full_matrices=True)

U, d, Vt = la.svd( mat, full_matrices=False )

assert np.all( d[:-1] >= d[1:] ) # sorted

eigen = d**2/2000

sumvariance = np.cumsum(eigen)

sumvariance /= sumvariance[-1]

i= np.arange(1,2001)

```

## Part 4: Get best value of 'k', based on explained variance



## of matrix U

```
[ ]: import matplotlib.pyplot as plt
      plt.plot( i, sumvariance)
      plt.rcParams["figure.figsize"] = [7,7]
```

## Part 5: TruncatedSVD on U to reduce U to 'k' components

```
[ ]: import numpy as np
      from sklearn.decomposition import TruncatedSVD

      component_matrix = []
      variance_matrix = []

      model = TruncatedSVD(n_components=250).fit(U)
      X_proj = model.transform(U)

      explained_variances = round(np.mean(np.var(X_proj, axis=0) / np.var(U, axis=0))
```

```
[ ]: from sklearn.decomposition import TruncatedSVD

      from scipy.sparse import csr_matrix

      standardized_data_sparse_train = csr_matrix(U)

      tsvd = TruncatedSVD(n_components=250)

      standardized_data_sparse_tsvd_train = tsvd.fit(standardized_data_sparse_train)
```

## Part 6: # Aggregate the Features and visualize the

```
[ ]:

from sklearn.cluster import KMeans
from sklearn.metrics import pairwise_distances_argmin_min
from scipy.spatial import distance
from scipy.spatial.distance import cdist

clusters=range(1,40)
meandist=[]

meandist = []
Inertia_matrix = []
Assignment_matrix = []

for k in clusters:
    kmeanModel = KMeans(init='k-means++',n_clusters=k).fit(standardized_data_sparse_tsvd_train)
    kmeanModel.fit(standardized_data_sparse_tsvd_train)
    Assignment_matrix.append(kmeanModel.predict(standardized_data_sparse_tsvd_train))
    #meandist.append(sum(np.min(cdist(bow_train_vector, kmeanModel.cluster_centers_, metric='euclidean'), axis=1)))
    Inertia_matrix.append( kmeanModel.inertia_)
```

## Finding Optimal K from the elbow plot

```
[ ]:

# Finding Optimal K from the elbow plot

fig = plt.figure()
ax = fig.add_subplot(111)

# Set the figure width and heights
fig_size = plt.rcParams["figure.figsize"]
fig_size[0] = 12
fig_size[1] = 12
plt.rcParams["figure.figsize"] = fig_size

plt.plot(clusters, Inertia_matrix)
```

```

12/13 ax.plot(clusters,Inertia_matrix, marker='o', markersize=4,markeredgewidth=5, m

plt.grid(True)
plt.xlabel('Number of Clusters')
plt.ylabel('Inertia')
plt.show()

```

```

[ ]: # Using Optimal K get the Centroids,Labels and Inertia details

```

```

[ ]: from sklearn import cluster
centroids,labels,inertia = cluster.k_means(standardized_data_sparse_tsvd_train

```

## Part 7: Take a word and gather 100 words closer to it based on Cosine Similarity

- Here the word taken is :

```

[77]: top_tfidf[8]

```

```
'cafe'
```

```

[78]: from scipy import linalg, mat, dot

a = standardized_data_sparse_tsvd_train[8]

cosine_similarity_values=[]
cosine_similarity_index = []

for i in range(0,2000):
    c = abs(dot(a,standardized_data_sparse_tsvd_train[i].T)/linalg.norm(a)/lin
    cosine_similarity_index.append(i)

```

```
12/13: cosine_similarity_values.append(c)
```

```
[79]: c = np.column_stack((cosine_similarity_index,cosine_similarity_values))
```

```
[81]: datanew = c[c[:,1].argsort()[::-1]]  
#print(datanew)
```

```
[82]: top100_words = (np.asarray(datanew)[0:100,])  
#print(top100_words)
```

```
[83]: top100_index= [int(row[0]) for row in top100_words]
```

```
[ ]: #np.asarray(top100_index).shape
```

```
[57]: top_closest_words= np.asarray(top_tfidf)[top100_index]
```

```
[84]: from sklearn.cluster import KMeans  
from sklearn.metrics import pairwise_distances_argmin_min  
from scipy.spatial import distance  
from scipy.spatial.distance import cdist  
  
clusters=range(1,5)  
meandist=[]  
  
meandist = []  
Inertia_matrix = []  
Assignment_matrix = []  
  
for k in clusters:  
    kmeanModel = KMeans(init='k-means++',n_clusters=k).fit(standardized_data_sparse_tsvd_train)  
    kmeanModel.fit(standardized_data_sparse_tsvd_train[top100_index])  
    Assignment_matrix.append(kmeanModel.predict(standardized_data_sparse_tsvd_train[top100_index]))  
    Inertia_matrix.append( kmeanModel.inertia_)
```

12/13

```
# Finding Optimal K from the elbow plot
```


```
fig = plt.figure()
ax = fig.add_subplot(111)
```

```
# Set the figure width and heights
```

```
fig_size = plt.rcParams["figure.figsize"]
fig_size[0] = 1
fig_size[1] = 1
plt.rcParams["figure.figsize"] = fig_size
```

```
plt.plot(clusters, Inertia_matrix)
ax.plot(clusters, Inertia_matrix, marker='o', markersize=4, markeredgewidth=5, ma
```

```
plt.grid(True)
plt.xlabel('Number of Clusters')
plt.ylabel('Inertia')
plt.show()
```



[87]:

```
from sklearn import cluster
centroids, labels, inertia = cluster.k_means(standardized_data_sparse_tsvd_train

# All the Cluster details:
{i: np.where(kmeanModel.labels_ == i)[0] for i in range(kmeanModel.n_clusters)}
```

```
{0: array([12, 69, 81, 86, 88]),
```

```
1: array([ 3,  4,  5,  7, 16, 20, 21, 22, 23, 24, 32, 35, 36, 39, 42, 45, 5
2,
          60, 61, 62, 67, 68, 70, 72, 75, 77, 91, 92, 98]),
2: array([ 0,  1,  6, 10, 11, 13, 15, 17, 18, 19, 25, 26, 27, 28, 29, 30, 3
1,
          34, 38, 40, 43, 46, 47, 49, 51, 53, 54, 55, 56, 57, 59, 63, 65, 66,
          71, 73, 74, 76, 78, 80, 82, 83, 85, 87, 90, 93, 94, 95, 96, 97, 99]),
3: array([ 2,  8,  9, 14, 33, 37, 41, 44, 48, 50, 58, 64, 79, 84, 89])}
```

```
[88]: # No of words in Cluster 0
      sum(kmeanModel.labels_ == 0)
```

```
5
```

```
[89]: # No of words in Cluster 1
      sum(kmeanModel.labels_ == 1)
```

```
29
```

```
[90]: # No of words in Cluster 2
      sum(kmeanModel.labels_ == 2)
```

```
51
```

```
[91]: # No of words in Cluster 3
      sum(kmeanModel.labels_ == 3)
```

```
15
```

## Part 8: Display the Clusters

```
[92]: #from sklearn import cluster
https://
```

12/13

```

#centroids,labels,inertia = cluster.k_means(standardized_data_sparse_tsvd_train

# All the Cluster details:
#{i: np.where(kmeanModel.labels_ == i)[0] for i in range(kmeanModel.n_clusters,

# Getting details of Cluster 5

index= np.where(kmeanModel.labels_ == 2)[0]
index_new = np.asarray(index)
cluster5 = str(np.asarray(X_train)[index_new,])

from os import path
from PIL import Image
from wordcloud import WordCloud, STOPWORDS, ImageColorGenerator

print("The WordCloud for Cluster 1 is given Below:")


wordcloud = WordCloud().generate(str(cluster5))

fig_size = plt.rcParams["figure.figsize"]
fig_size[0] = 20
fig_size[1] = 20
plt.rcParams["figure.figsize"] = fig_size

plt.imshow(wordcloud, interpolation='bilinear')
plt.show()

```

The WordCloud for Cluster 1 is given Below:



This cluster is themed around "Cafe"- we have words like coffe, tea, cup, breakfast, cereal,espresso,eat,chai aggregated around this theme.

## MODULE 2: ANALYSIS OF TOP 5000 TFIDF FEATURES

## # Part 1: get the Top 5k TFIDF Features

```
[93]: # source: https://buhrmann.github.io/tfidf-analysis.html
def top_tfidf_feats(row, features, top_n=5000):
    ''' Get top n tfidf values in row and return them with their corresponding
    topn_ids = np.argsort(row)[::-1][:top_n]
    #top_feats = [(features[i], row[i]) for i in topn_ids]
    top_feats = [(features[i]) for i in topn_ids]
    df = pd.DataFrame(top_feats)
    #df.columns = ['feature', 'tfidf']
    df.columns = ['feature']
    return top_feats
```

https://



```
12/13 top_tfidf = top_tfidf_feats(tfidf_train[1,:].toarray()[0],features,5000)
```

## Part 2 : Created Word Co-Occurrence Matrix with neighbourhood = 5

```
[94]: import numpy as np
import pandas as pd

ctxs = list(top_tfidf)

l_unique = list(set((' '.join(ctxs)).split(' ')))
mat = np.zeros((len(l_unique), len(l_unique)))

nei = []
nei_size = 5

for ctx in ctxs:
    words = ctx.split(' ')

    for i, _ in enumerate(words):
        nei.append(words[i])

    if len(nei) > (nei_size * 2) + 1:
        nei.pop(0)

    pos = int(len(nei) / 2)
    for j, _ in enumerate(nei):
        if nei[j] in l_unique and words[i] in l_unique:
            mat[l_unique.index(nei[j]), l_unique.index(words[i])] += 1

mat = pd.DataFrame(mat)
mat.index = l_unique
mat.columns = l_unique
print(mat)
```

	drunk	hunger	evo	hondura	...	boil	epa	bittersweet	smell
drunk	1.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
hunger	0.0	1.0	0.0	0.0	...	0.0	0.0	0.0	0.0
evo	0.0	0.0	1.0	0.0	...	0.0	0.0	0.0	0.0
hondura	0.0	0.0	0.0	1.0	...	0.0	0.0	0.0	0.0
box	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0

fret	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
ginseng	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
hcg	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
impuls	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
accomod	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
help	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
farley	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
hors	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
boscoli	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
hazan	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
heath	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
dougla	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
hectic	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
amazonprim	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
jake	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
domino	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
landmin	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
kenya	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
hlaf	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
boba	0.0	0.0	0.0	0.0	...	1.0	0.0	0.0	0.0
genet	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
appletini	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
evidenc	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
harlan	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
gem	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
...	...	...	...	...	...	...	...	...	...
barney	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
guzzl	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
frere	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
konbu	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
havoc	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
johnson	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
crop	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
growl	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
heavier	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
elev	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
bald	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
equilibrium	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
famlili	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
afer	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
assess	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
bind	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0

earthborn	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
boquet	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
crummi	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
enterpris	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
lamb	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
henri	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
joyva	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
led	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
headach	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
krusteaz	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0
boil	0.0	0.0	0.0	0.0	...	1.0	0.0	0.0	0.0
epa	0.0	0.0	0.0	0.0	...	0.0	1.0	0.0	0.0
bittersweet	0.0	0.0	0.0	0.0	...	0.0	0.0	1.0	0.0
smell	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	1.0

[5000 rows x 5000 columns]

## ## Part 3: Word Co-Occurrence Matrix Decomposition using SVD

```

12/13/ import numpy as np

la= np.linalg

#u, s, vh = la.svd(mat, full_matrices=True)

U, d, Vt = la.svd( mat, full_matrices=False )

assert np.all( d[:-1] >= d[1:] ) # sorted

eigen = d**2/5000

sumvariance = np.cumsum(eigen)

sumvariance /= sumvariance[-1]

i= np.arange(1,5001)


```

Part 4: Get best value of 'k', based on explained variance of matrix U

```

[97]: import matplotlib.pyplot as plt
plt.plot( i,sumvariance)
plt.rcParams["figure.figsize"]= [2,2]

```



Part 5: TruncatedSVD on U to reduce U to 'k' components (from the elbow curve)

```
[98]: import numpy as np
from sklearn.decomposition import TruncatedSVD

component_matrix = []
variance_matrix = []

model = TruncatedSVD(n_components=500).fit(U)
X_proj = model.transform(U)

explained_variances = round(np.mean(np.var(X_proj, axis=0) / np.var(U, axis=0))
```

```
[103]: from sklearn.decomposition import TruncatedSVD

from scipy.sparse import csr_matrix

standardized_data_sparse_train = csr_matrix(U)
```

<https://>

12/13

```
tsvd = TruncatedSVD(n_components=500)

standardized_data_sparse_tsvd_train = tsvd.fit(standardized_data_sparse_train)
```

## Part 6: # Aggregate the Features and visualize the Clusters (K Means)

```
[100]: from sklearn.cluster import KMeans
        from sklearn.metrics import pairwise_distances_argmin_min
        from scipy.spatial import distance
        from scipy.spatial.distance import cdist

        clusters=range(1,40)
        meandist=[]

        meandist = []
        Inertia_matrix = []
        Assignment_matrix = []

        for k in clusters:
            kmeanModel = KMeans(init='k-means++',n_clusters=k).fit(standardized_data_sparse_tsvd_train)
            kmeanModel.fit(standardized_data_sparse_tsvd_train)
            Assignment_matrix.append(kmeanModel.predict(standardized_data_sparse_tsvd_train))
            #meandist.append(sum(np.min(cdist(bow_train_vector, kmeanModel.cluster_centers_, metric='euclidean'), axis=1)))
            Inertia_matrix.append(kmeanModel.inertia_)
```

## Finding Optimal K from the elbow plot

```
[105]: # Finding Optimal K from the elbow plot

        fig = plt.figure()
        ax = fig.add_subplot(111)

        # Set the figure width and heights
        fig_size = plt.rcParams["figure.figsize"]
        fig_size[0] = 20
```


https://

12/13/

```
fig_size[1] = 20
plt.rcParams["figure.figsize"] = fig_size

plt.plot(clusters, Inertia_matrix)
ax.plot(clusters, Inertia_matrix, marker='o', markersize=4, markeredgewidth=5, ma

plt.grid(True)
plt.xlabel('Number of Clusters')
plt.ylabel('Inertia')
plt.show()
```





```
[106]: # Using Optimal K get the Centroids, Labels and Inertia details
from sklearn import cluster
centroids, labels, inertia = cluster.k_means(standardized_data_sparse_tsvd_train
```

## Part 7: Take a word and gather 100 words closer to it based on Cosine Similarity

### Here the word taken is

```
[157]: top_tfidf[100]
```

```
'eighteen'
```

```
[158]: from scipy import linalg, mat, dot

a = standardized_data_sparse_tsvd_train[100]

cosine_similarity_values=[]
cosine_similarity_index = []

for i in range(0,5000):
    c = abs(dot(a,standardized_data_sparse_tsvd_train[i].T)/linalg.norm(a)/linalg.norm(standardized_data_sparse_tsvd_train[i]))
    cosine_similarity_index.append(i)
```

```
[159]: c = np.column_stack((cosine_similarity_index,cosine_similarity_values))

datanew = c[c[:,1].argsort()[::-1]]

top100_words = (np.asarray(datanew)[0:100,])

top100_index= [int(row[0]) for row in top100_words]

top_closest_words= np.asarray(top_tfidf)[top100_index]
```

```
[160]: from sklearn.cluster import KMeans
from sklearn.metrics import pairwise_distances_argmin_min
from scipy.spatial import distance
from scipy.spatial.distance import cdist

clusters=range(1,5)
meandist=[]

meandist = []
Inertia_matrix = []
Assignment_matrix = []

for k in clusters:
    kmeanModel = KMeans(init='k-means++',n_clusters=k).fit(standardized_data_sparse_tsvd_train)
    kmeanModel.fit(standardized_data_sparse_tsvd_train[top100_index])
    Assignment_matrix.append(kmeanModel.predict(standardized_data_sparse_tsvd_train[top100_index]))
    Inertia_matrix.append( kmeanModel.inertia_)
```

```
[162]: # Finding Optimal K from the elbow plot

fig = plt.figure()
#ax = fig.add_subplot(111)

# Set the figure width and heights
```


```

12/13/ fig_size = plt.rcParams["figure.figsize"]
fig_size[0] = 3
fig_size[1] = 3
plt.rcParams["figure.figsize"] = fig_size

plt.plot(clusters, Inertia_matrix)
ax.plot(clusters, Inertia_matrix, marker='o', markersize=4, markeredgewidth=5, m

plt.grid(True)
plt.xlabel('Number of Clusters')
plt.ylabel('Inertia')

```



```

[163]: from sklearn import cluster
centroids, labels, inertia = cluster.k_means(standardized_data_sparse_tsvd_train

# All the Cluster details:
{i: np.where(kmeanModel.labels_ == i)[0] for i in range(kmeanModel.n_clusters)}

{0: array([ 3, 15, 48, 58, 74, 86, 90]),
 1: array([ 1,  6,  8, 14, 24, 28, 30, 31, 33, 46, 52, 53, 54, 55, 60, 62, 64,
          68, 69, 70, 71, 76, 77, 78, 80, 82, 89, 95, 96]),
 2: array([ 7, 16, 22, 25, 29, 34, 43, 50, 56, 75, 98]),
 3: array([ 0,  2,  4,  5,  9, 10, 11, 12, 13, 17, 18, 19, 20, 21, 23, 26, 27,
          32, 35, 36, 37, 38, 39, 40, 41, 42, 44, 45, 47, 49, 51, 57, 59, 61,
          63, 65, 66, 67, 72, 73, 79, 81, 83, 84, 85, 87, 88, 91, 92, 93, 94,
          97, 99])}

```

```

[168]: # No of words in Cluster 0

```

https://

```
12/13 sum(kmeanModel.labels_ == 0)
```

```
7
```

```
[167]: # No of words in Cluster 1
sum(kmeanModel.labels_ == 1)
```

```
29
```

```
[166]: # No of words in Cluster 2
sum(kmeanModel.labels_ == 2)
```

```
11
```

```
[165]: # No of words in Cluster 3
sum(kmeanModel.labels_ == 3)
```

```
53
```

---

## # Part 8: Display the Clusters with max number of words

```
[169]: index= np.where(kmeanModel.labels_ == 3)[0]
index_new = np.asarray(index)
cluster5 = str(np.asarray(X_train)[index_new,])

from os import path
from PIL import Image
from wordcloud import WordCloud, STOPWORDS, ImageColorGenerator

print("The WordCloud for the Cluster is given Below:")
```

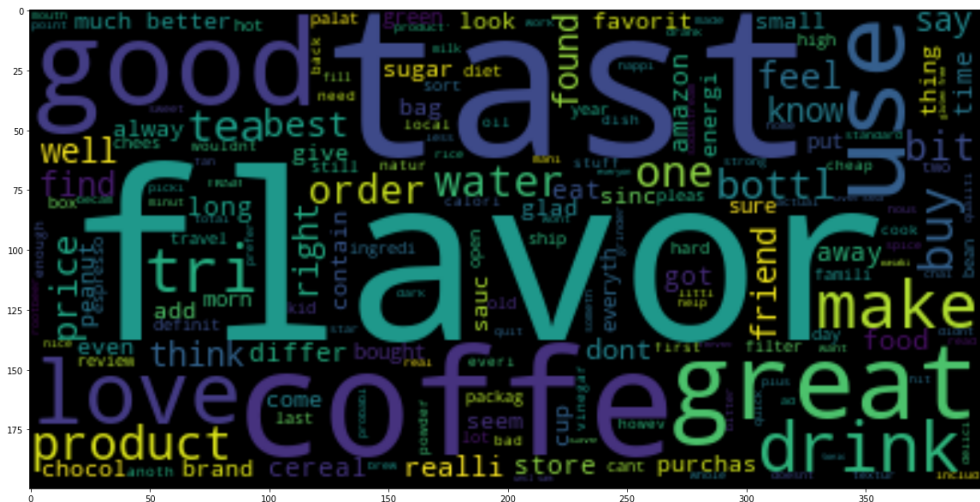
https://

```
wordcloud = WordCloud().generate(str(cluster5))

fig_size = plt.rcParams["figure.figsize"]
fig_size[0] = 20
fig_size[1] = 20
plt.rcParams["figure.figsize"] = fig_size

plt.imshow(wordcloud, interpolation='bilinear')
plt.show()
```

The WordCloud for the Cluster is given Below:



Cluster theme is "Diet" Indicated by words like cereal, energi, Green, diet etc.

## Module 3: Analysis of top 10K TFIDF features

## Part 1: get the Top 10k TFIDF Features

```
[12/13] # source: https://buhrmann.github.io/tfidf-analysis.html
def top_tfidf_feats(row, features, top_n=10000):
    ''' Get top n tfidf values in row and return them with their corresponding
    topn_ids = np.argsort(row)[::-1][:top_n]
    #top_feats = [(features[i], row[i]) for i in topn_ids]
    top_feats = [(features[i]) for i in topn_ids]
    df = pd.DataFrame(top_feats)
    #df.columns = ['feature', 'tfidf']
    df.columns = ['feature']
    return top_feats

top_tfidf = top_tfidf_feats(tfidf_train[1,:].toarray()[0], features, 10000)
```

Part 2 : Created Word Co-Occurrence Matrix with neighbourhood = 5

```
[189]: import numpy as np
import pandas as pd

ctxs = list(top_tfidf)

l_unique = list(set((' '.join(ctxs)).split(' ')))
mat = np.zeros((len(l_unique), len(l_unique)))

nei = []
nei_size = 5

for ctx in ctxs:
    words = ctx.split(' ')

    for i, _ in enumerate(words):
        nei.append(words[i])

    if len(nei) > (nei_size * 2) + 1:
        nei.pop(0)

    pos = int(len(nei) / 2)
    for j, _ in enumerate(nei):
        if nei[j] in l_unique and words[i] in l_unique:
            mat[l_unique.index(nei[j]), l_unique.index(words[i])] += 1
```

12/13

```
mat = pd.DataFrame(mat)
mat.index = l_unique
mat.columns = l_unique
print(mat)
```

	underneath	hondura	constitu	...	lamb	headach	joyva
underneath	1.0	0.0	0.0	...	0.0	0.0	0.0
hondura	0.0	1.0	0.0	...	0.0	0.0	0.0
constitu	0.0	0.0	1.0	...	0.0	0.0	0.0
copycat	0.0	0.0	0.0	...	0.0	0.0	0.0
ginseng	0.0	0.0	0.0	...	0.0	0.0	0.0
take	0.0	0.0	0.0	...	0.0	0.0	0.0
impuls	0.0	0.0	0.0	...	0.0	0.0	0.0
tastybit	0.0	0.0	0.0	...	0.0	0.0	0.0
help	0.0	0.0	0.0	...	0.0	0.0	0.0
hectic	0.0	0.0	0.0	...	0.0	0.0	0.0
jake	0.0	0.0	0.0	...	0.0	0.0	0.0
hlaf	0.0	0.0	0.0	...	0.0	0.0	0.0
kenya	0.0	0.0	0.0	...	0.0	0.0	0.0
theater	0.0	0.0	0.0	...	0.0	0.0	0.0
snot	0.0	0.0	0.0	...	0.0	0.0	0.0
splinter	0.0	0.0	0.0	...	0.0	0.0	0.0
stubborn	0.0	0.0	0.0	...	0.0	0.0	0.0
harlan	0.0	0.0	0.0	...	0.0	0.0	0.0
scrape	0.0	0.0	0.0	...	0.0	0.0	0.0
wast	0.0	0.0	0.0	...	0.0	0.0	0.0
stage	0.0	0.0	0.0	...	0.0	0.0	0.0
col	0.0	0.0	0.0	...	0.0	0.0	0.0
falafel	0.0	0.0	0.0	...	0.0	0.0	0.0
dramat	0.0	0.0	0.0	...	0.0	0.0	0.0
art	0.0	0.0	0.0	...	0.0	0.0	0.0
harbor	0.0	0.0	0.0	...	0.0	0.0	0.0
spasm	0.0	0.0	0.0	...	0.0	0.0	0.0
duke	0.0	0.0	0.0	...	0.0	0.0	0.0
jazz	0.0	0.0	0.0	...	0.0	0.0	0.0
thin	0.0	0.0	0.0	...	0.0	0.0	0.0
...	...	...	...	...	...	...	...
gourd	0.0	0.0	0.0	...	0.0	0.0	0.0
swoon	0.0	0.0	0.0	...	0.0	0.0	0.0
darwin	0.0	0.0	0.0	...	0.0	0.0	0.0
dosent	0.0	0.0	0.0	...	0.0	0.0	0.0
tealik	0.0	0.0	0.0	...	0.0	0.0	0.0

funnel	0.0	0.0	0.0	...	0.0	0.0	0.0
jewel	0.0	0.0	0.0	...	0.0	0.0	0.0
nug	0.0	0.0	0.0	...	0.0	0.0	0.0
airless	0.0	0.0	0.0	...	0.0	0.0	0.0
trap	0.0	0.0	0.0	...	0.0	0.0	0.0
doodl	0.0	0.0	0.0	...	0.0	0.0	0.0
held	0.0	0.0	0.0	...	0.0	0.0	0.0
gobstopp	0.0	0.0	0.0	...	0.0	0.0	0.0
goop	0.0	0.0	0.0	...	0.0	0.0	0.0
tidi	0.0	0.0	0.0	...	0.0	0.0	0.0
nostril	0.0	0.0	0.0	...	0.0	0.0	0.0
guzzl	0.0	0.0	0.0	...	0.0	0.0	0.0
upright	0.0	0.0	0.0	...	0.0	0.0	0.0
unison	0.0	0.0	0.0	...	0.0	0.0	0.0
voxbox	0.0	0.0	0.0	...	0.0	0.0	0.0
equilibrium	0.0	0.0	0.0	...	0.0	0.0	0.0
shift	0.0	0.0	0.0	...	0.0	0.0	0.0
unapp	0.0	0.0	0.0	...	0.0	0.0	0.0
vanish	0.0	0.0	0.0	...	0.0	0.0	0.0
chiavetta	0.0	0.0	0.0	...	0.0	0.0	0.0
creek	0.0	0.0	0.0	...	0.0	0.0	0.0
henri	0.0	0.0	0.0	...	0.0	0.0	0.0
lamb	0.0	0.0	0.0	...	1.0	0.0	0.0
headach	0.0	0.0	0.0	...	0.0	1.0	0.0
joyva	0.0	0.0	0.0	...	0.0	0.0	1.0

[10000 rows x 10000 columns]

### Part 3: Word Co-Occurrence Matrix Decomposition using SVD

```
[190]: import numpy as np

la= np.linalg

#u, s, vh = la.svd(mat, full_matrices=True)

U, d, Vt = la.svd( mat, full_matrices=False )

assert np.all( d[:-1] >= d[1:] ) # sorted
```

<https://>



12/13/


```
eigen = d**2/10000

sumvariance = np.cumsum(eigen)

sumvariance /= sumvariance[-1]
```

Part 4: Get best value of 'k', based on explained variance of matrix U

```
[192]: import matplotlib.pyplot as plt
plt.plot( i,sumvariance)
plt.rcParams["figure.figsize"]= [2,2]
```



Part 5: TruncatedSVD on U to reduce U to 'k' components

```
[193]: import numpy as np
```

<https://>

```

12/13 from sklearn.decomposition import TruncatedSVD

component_matrix = []
variance_matrix = []

model = TruncatedSVD(n_components=1200).fit(U)
X_proj = model.transform(U)

explained_variances = round(np.mean(np.var(X_proj, axis=0)) / np.var(U, axis=0))

```

```

[194]: from sklearn.decomposition import TruncatedSVD

from scipy.sparse import csr_matrix

standardized_data_sparse_train = csr_matrix(U)

tsvd = TruncatedSVD(n_components=1200)

standardized_data_sparse_tsvd_train = tsvd.fit(standardized_data_sparse_train)

```

## Part 6: # Aggregate the Features and visualize the Clusters (K Means)

```

[195]: from sklearn.cluster import KMeans
from sklearn.metrics import pairwise_distances_argmin_min
from scipy.spatial import distance
from scipy.spatial.distance import cdist

clusters=range(1,40)
meandist=[]

meandist = []
Inertia_matrix = []
Assignment_matrix = []

for k in clusters:
    kmeanModel = KMeans(init='k-means++', n_clusters=k).fit(standardized_data_sparse_train)
    kmeanModel.fit(standardized_data_sparse_tsvd_train)
    Assignment_matrix.append(kmeanModel.predict(standardized_data_sparse_tsvd_train))

```

## Finding Optimal K from the elbow plot


```
[196]: # Finding Optimal K from the elbow plot

fig = plt.figure()
ax = fig.add_subplot(111)

# Set the figure width and heights
fig_size = plt.rcParams["figure.figsize"]
fig_size[0] = 12
fig_size[1] = 12
plt.rcParams["figure.figsize"] = fig_size

plt.plot(clusters, Inertia_matrix)
ax.plot(clusters, Inertia_matrix, marker='o', markersize=4, markeredgewidth=5, m

plt.grid(True)
plt.xlabel('Number of Clusters')
plt.ylabel('Inertia')
plt.show()
```



```
[197]: from sklearn import cluster
centroids, labels, inertia = cluster.k_means(standardized_data_sparse_tsvd_train
```

## Part 7: Take a word and gather 100 words closer to it based on Cosine Similarity

Here the word taken is :

[219]:

```
top_tfidf[1234]
```

```
'enfamil'
```

[220]:

```
from scipy import linalg, mat, dot

a = standardized_data_sparse_tsvd_train[1234]

cosine_similarity_values=[]
cosine_similarity_index = []

for i in range(0,10000):
    c = abs(dot(a,standardized_data_sparse_tsvd_train[i].T)/linalg.norm(a)/linalg.norm(standardized_data_sparse_tsvd_train[i]))
    cosine_similarity_index.append(i)
    cosine_similarity_values.append(c)
```

[221]:

```
c = np.column_stack((cosine_similarity_index,cosine_similarity_values))

datanew = c[c[:,1].argsort()[::-1]]

top100_words = (np.asarray(datanew)[0:100,])

top100_index= [int(row[0]) for row in top100_words]

top_closest_words= np.asarray(top_tfidf)[top100_index]
```

[12/13]

```
from sklearn.cluster import KMeans
from sklearn.metrics import pairwise_distances_argmin_min
from scipy.spatial import distance
from scipy.spatial.distance import cdist

clusters=range(1,6)
meandist=[]

meandist = []
Inertia_matrix = []
Assignment_matrix = []

for k in clusters:
    kmeanModel = KMeans(init='k-means++',n_clusters=k).fit(standardized_data_sparse_tsvd_train[0:k])
    kmeanModel.fit(standardized_data_sparse_tsvd_train[top100_index])
    Assignment_matrix.append(kmeanModel.predict(standardized_data_sparse_tsvd_train[top100_index]))
    Inertia_matrix.append( kmeanModel.inertia_)
```

[223]:


```
# Finding Optimal K from the elbow plot

fig = plt.figure()
ax = fig.add_subplot(111)

# Set the figure width and heights
fig_size = plt.rcParams["figure.figsize"]
fig_size[0] = 3
fig_size[1] = 3
plt.rcParams["figure.figsize"] = fig_size

plt.plot(clusters, Inertia_matrix)
ax.plot(clusters,Inertia_matrix, marker='o', markersize=4,markeredgewidth=5, markeredgecolor='r')

plt.grid(True)
plt.xlabel('Number of Clusters')
plt.ylabel('Inertia')
plt.show()
```



```
[224]: from sklearn import cluster
centroids, labels, inertia = cluster.k_means(standardized_data_sparse_tsvd_train

# All the Cluster details:

{i: np.where(kmeanModel.labels_ == i)[0] for i in range(kmeanModel.n_clusters)}
```

```
{0: array([ 3, 10, 21, 33, 40, 69, 74, 79, 92, 96]),
 1: array([ 4,  9, 11, 16, 20, 23, 24, 34, 36, 43, 44, 45, 46, 47, 50, 51, 55,
          56, 57, 58, 61, 65, 66, 77, 83, 84, 86, 95, 98, 99]),
 2: array([ 0,  5,  6,  7,  8, 12, 13, 14, 15, 18, 19, 26, 27, 28, 29, 31, 32,
          35, 39, 41, 49, 54, 59, 62, 67, 68, 70, 72, 73, 75, 78, 82, 87, 88,
          91, 93, 97]),
 3: array([17, 42]),
 4: array([ 1,  2, 22, 25, 30, 37, 38, 48, 52, 53, 60, 63, 64, 71, 76, 80, 81,
          85, 89, 90, 94])}
```

```
[225]: # No of words in Cluster 1
sum(kmeanModel.labels_ == 0)
```

```
10
```

```
[226]: # No of words in Cluster 2
sum(kmeanModel.labels_ == 1)
```

```
30
```

```
[227]: # No of words in Cluster 3
sum(kmeanModel.labels_ == 2)
```

```
37
```

```
[228]: # No of words in Cluster 4
sum(kmeanModel.labels_ == 3)
```

2

## Part 8: Display the Cluster which has the maximum number of words

```
[230]: index= np.where(kmeanModel.labels_ == 2)[0]
index_new = np.asarray(index)
cluster5 = str(np.asarray(X_train)[index_new,])

from os import path
from PIL import Image
from wordcloud import WordCloud, STOPWORDS, ImageColorGenerator

print("The WordCloud for the cluster is given Below:")

wordcloud = WordCloud().generate(str(cluster5))

fig_size = plt.rcParams["figure.figsize"]
fig_size[0] = 25
fig_size[1] = 25
plt.rcParams["figure.figsize"] = fig_size

plt.imshow(wordcloud, interpolation='bilinear')
plt.show()
```

The WordCloud for the cluster is given Below:



